

## HW 4

```
In [1]: # Eric Liu  
        # Eric Wang  
        # Austin Du
```

### Imports

```
In [2]: library("ggplot2")  
        library("gridExtra")  
        library("dplyr")  
        library("reshape2")  
        library("visreg")
```

Attaching package: 'dplyr'

The following object is masked from 'package:gridExtra':

combine

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

### Load Data

```
In [3]: gauge = read.table("./gauge.txt", sep=" ", header=T)
density = gauge$density
gain = gauge$gain
head(gauge)
```

Out[3]: A data.frame: 6 × 2

	density	gain
	<dbl>	<dbl>
1	0.686	17.6
2	0.686	17.3
3	0.686	16.9
4	0.686	16.2
5	0.686	17.1
6	0.686	18.5

```
In [4]: # the unique densities of the blocks
sort(unique(gauge$density))
```

Out[4]: 0.001 · 0.08 · 0.148 · 0.223 · 0.318 · 0.412 · 0.508 · 0.604 · 0.686

## Best Fit Line

```
In [5]: # function to print out the equation of the line and the sum of squared residuals
summarize_fit = function(line) {
  cat("intercept:\t\t", line$coefficients[1], "\n")
  cat("slope:\t\t\t", line$coefficients[2], "\n")
  cat("sum squared residuals:\t", sum(resid(line)^2), "\n")
  cat("r squared:\t\t", summary(line)$r.squared)
}
```

```
In [6]: # without any transformations, fit the data with linear regression and plot the residuals
reg_plot = ggplot(gauge, aes(x=density, y=gain)) +
  geom_point() +
  stat_smooth(method='lm', formula = y~x) +
  labs(title="Density vs Gain", x="Density", y="Gain")

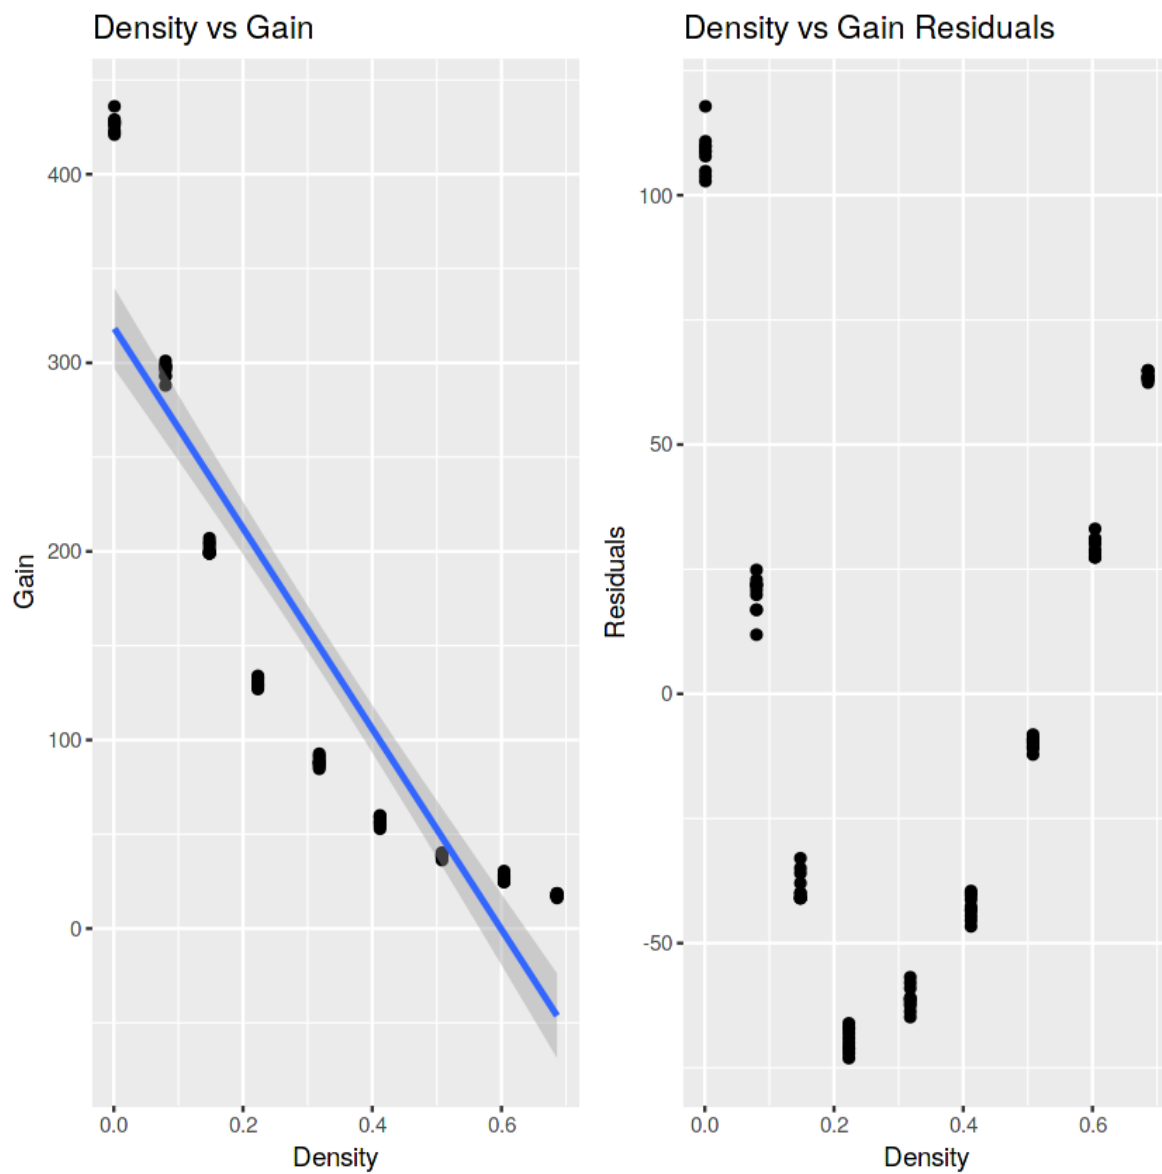
# compute the residuals for the linear regression and plot
lin_reg = lm(gauge$gain ~ gauge$density)
residuals = data.frame(x=gauge$density, resid=resid(lin_reg))
resid_plot = ggplot(residuals, aes(x=x, y=resid)) +
  geom_point() +
  labs(title="Density vs Gain Residuals", x="Density", y="Residuals")

# put both plots side by side
grid.arrange(reg_plot, resid_plot, ncol=2)

# print out the equation of the line and the sum of squared residuals
summarize_fit(lin_reg)
```

```
intercept:      318.7015  
slope:         -531.9507  
sum squared residuals: 291335  
r squared:      0.8156974
```

Out[6]:



Since the residuals show a polynomial pattern, try doing polynomial regression

```
In [7]: # transform data (take the sqrt of y axis) and fit to linear regression
gauge$sqrt_gain = sqrt(gauge$gain)
reg_plot = ggplot(gauge, aes(x=density, y=sqrt_gain)) +
  geom_point() +
  stat_smooth(method='lm', formula = y~x) +
  labs(title="Density vs Sqrt. Gain", x="Density", y="Square Root of Ga
in")


# compute the residuals for the linear regression and plot
lin_reg = lm(gauge$sqrt_gain ~ gauge$density)
residuals = data.frame(x=gauge$density, resid=resid(lin_reg))
resid_plot = ggplot(residuals, aes(x=x, y=resid)) +
  geom_point() +
  labs(title="Density vs Sqrt. Gain Residuals", x="Density", y="Residua
ls")

# put both plots side by side
grid.arrange(reg_plot, resid_plot, ncol=2)

# print out the equation of the line and the sum of squared residuals
summarize_fit(lin_reg)
```

```
intercept:      18.28919  
slope:         -23.02002  
sum squared residuals: 175.6669  
r squared:      0.9321839
```

Out[7]:



Since the residuals still show a polynomial pattern, try cubic regression

```
In [8]: # transform data (take the cube root of y axis) and fit to linear regression
gauche$cubert_gain = (gauche$gain)^(1/3)
reg_plot = ggplot(gauche, aes(x=density, y=cubert_gain)) +
  geom_point() +
  stat_smooth(method='lm', formula = y~x) +
  labs(title="Density vs Cbrt. of Gain", x="Density", y="Cube Root of Gain")

# compute the residuals for the linear regression and plot
lin_reg = lm(gauche$cubert_gain ~ gauche$density)
residuals = data.frame(x=gauche$density, resid=resid(lin_reg))
resid_plot = ggplot(residuals, aes(x=x, y=resid)) +
  geom_point() +
  labs(title="Density vs Cbrt. of Gain Residuals", x="Density", y="Residuals")


# put both plots side by side
grid.arrange(reg_plot, resid_plot, ncol=2)

# print out the equation of the line and the sum of squared residuals
summarize_fit(lin_reg)
```


```
intercept:      7.03021  
slope:         -7.004705  
sum squared residuals: 8.76708  
r squared:     0.9622668
```

Out[8]:

Density vs Cbrt. of Gain



Density vs Cbrt. of Gain Residuals



Since the residuals still show a polynomial residuals graph, let's try exponential distribution



```
In [9]: # transform data (take the ln of y axis) and fit to linear regression
gauge$ln_gain = log(gauge$gain)
reg_plot = ggplot(gauge, aes(x=density, y=ln_gain)) +
  geom_point() +
  stat_smooth(method='lm', formula = y~x) +
  labs(title="Density vs ln Gain", x="Density", y="Natural Log of Gain"
)

# compute the residuals for the linear regression and plot
lin_reg = lm(gauge$ln_gain ~ gauge$density)
residuals = data.frame(x=gauge$density, resid=resid(lin_reg))
resid_plot = ggplot(residuals, aes(x=x, y=resid)) +
  geom_point() +
  labs(title="Density vs ln Gain Residuals", x="Density", y="Residuals"
)


# put both plots side by side
grid.arrange(reg_plot, resid_plot, ncol=2)

# print out the equation of the line and the sum of squared residuals
summarize_fit(lin_reg)
```


```
intercept:      5.997265  
slope:         -4.605937  
sum squared residuals: 0.4059345  
r squared:      0.9958183
```

Out[9]:

Density vs ln Gain



Density vs ln Gain Residuals



**Below is the model we propose**

```
In [10]: # transform data (take the sqrt of the ln of y axis) and fit to linear re
         gression
         gauge$sqrt_ln_gain = sqrt(log(gauge$gain))
         reg_plot = ggplot(gauge, aes(x=density, y=sqrt_ln_gain)) +
           geom_point() +
           stat_smooth(method='lm', formula = y~x) +
           labs(title="Density vs Sqrt. ln Gain", x="Density", y="Square Root of
         Natural Log of Gain")

         # compute the residuals for the linear regression and plot
         lin_reg = lm(gauge$sqrt_ln_gain ~ gauge$density)
         residuals = data.frame(x=gauge$density, resid=resid(lin_reg))
         resid_plot = ggplot(residuals, aes(x=x, y=resid)) +
           geom_point() +
           labs(title="Density vs Sqrt. ln Gain Residuals", x="Density", y="Resi
         duals")


         # put both plots side by side
         grid.arrange(reg_plot, resid_plot, ncol=2)

         # print out the equation of the line and the sum of squared residuals
         summarize_fit(lin_reg)
```


```
intercept:      2.465989  
slope:         -1.105212  
sum squared residuals: 0.01539044  
r squared:      0.9972425
```

Out[10]:

Density vs Sqrt. ln Gain




Density vs Sqrt. ln Gain Residuals



```
In [11]: ggplot(residuals, aes(x=resid)) +
  geom_density() +
  labs(title="Density vs Sqrt. In Gain Residuals Density Plot", x="Residuals", y="Density")
```

Out[11]: **Density vs Sqrt. In Gain Residuals Density Plot**



```
In [0]: m = -1.105212
b = 2.465989

f = function(x) exp(1)^((m * x + b)^2)

# plot the original data but with our proposed regression line
ggplot(gauge, aes(x=density, y=gain)) +
  geom_point() +
  stat_function(fun=f, color="blue") +
  labs(title="Density vs Gain", x="Density", y="Gain") +
  annotate(geom="text", x=.35, y=200, label="gain = e^((-1.1052 * density + 2.4660)^2)", color="blue")
```

## What if data was not reported accurately?

My plan is to simulate a large number of times the following procedure:

1. For each density value, vary the value by  $\pm 10\%$  (or some other value in the range with a uniform distribution)
2. Fit the best line again and record the coefficients
3. Plot the distributions of the coefficients and analyze
4. ???
5. Profit

```
In [0]: # make a copy of the original df to not screw up anything later
# since we will always need to take the log of gain, do so now
original_df = data.frame(select(gauge, sqrt_ln_gain, density))
head(original_df)
```

```
In [0]: # parameters
N = 5000          # how many times to simulate
V = .15           # range of variation for density (from -V to +V)

# create lists to hold the simulated coefficients
slopes = vector("list", length=N)
intercepts = vector("list", length=N)

# create a function to vary the density column by any amount between -V and +V
vary_density = function(x) {
  vary_amount = runif(n=1, min=-V, max=V) + 1
  return(x * vary_amount)
}

# simulate N times
for(i in 1:N) {
  # first create a copy of the original df so our changes don't carry over between iterations
  copied_df = data.frame(original_df)

  # modify the density column with the random variation
  copied_df$density = unlist(lapply(copied_df$density, vary_density), use.names=F)

  # do linear regression as before and record the coefficients
  lin_reg = lm(sqrt_ln_gain ~ density, data=copied_df)
  slopes[i] = lin_reg$coefficients[2]
  intercepts[i] = lin_reg$coefficients[1]
}
```

```
In [0]: # WARNING: DO NOT UNCOMMENT AND RUN THIS CELL UNLESS YOU WANT TO HAVE TO
# REDO THE CUSTOM TEXT ON THE FOLLOWING GRAPHS

# # combine the lists into a single dataframe and save
# to_save = do.call(rbind, Map(data.frame, slope=slopes, intercept=interc
# epts))
# write.table(to_save, "simulated_density_coefficients.txt", sep=" ", ro
# w.names=F)
```

```
In [0]: # load the saved data
sdc = read.table("./simulated_density_coefficients.txt", sep="", header=T
)

# plot the distribution of slopes and intercepts
slope_plot = ggplot(sdc, aes(x=slope)) +
  geom_density() +
  geom_vline(aes(xintercept=mean(slope)), color="blue") +
  annotate("text", label=paste("mean slope at ", round(mean(sdc$slope),
digits=4)), x=mean(sdc$slope)+.005, y=10, angle=90) +
  geom_vline(aes(xintercept=m), color="red") +
  annotate("text", label=paste("actual slope at ", m), x=m+.005, y=5, a
ngle=90) +
  labs(title="Density Plot of Slope", x="Slope", y="Density")
intercept_plot = ggplot(sdc, aes(x=intercept)) +
  geom_density() +
  geom_vline(aes(xintercept=mean(intercept)), color="blue") +
  annotate("text", label=paste("mean intercept at ", round(mean(sdc$int
ercept), digits=4)), x=mean(sdc$intercept)+.001, y=50, angle=90) +
  geom_vline(aes(xintercept=b), color="red") +
  annotate("text", label=paste("actual intercept at ", b), x=b+.001, y=
50, angle=90) +
  labs(title="Density Plot of Intercept", x="Intercept", y="Density")

# put both plots side by side
grid.arrange(slope_plot, intercept_plot, ncol=2)
```

```
In [0]: # constant for percent variation
C = .01
```

```
In [0]: # get the mean and standard deviation of simulated slope
slope_mean = mean(sdc$slope)
slope_std = sd(sdc$slope)

# compute the probability of getting within C of the observed value
abs(pnorm(m + (m * C), mean=slope_mean, sd=slope_std, lower.tail=T) - pno
rm(m - (m * C), mean=slope_mean, sd=slope_std, lower.tail=T))
```

```
In [0]: # get the mean and standard deviation of simulated intercept
intercept_mean = mean(sdc$intercept)
intercept_std = sd(sdc$intercept)

# compute the probability of getting within C of the observed value
abs(pnorm(b + (b * C), mean=intercept_mean, sd=intercept_std, lower.tail=
T) - pnorm(b - (b * C), mean=intercept_mean, sd=intercept_std, lower.tail
=T))
```

## Predicting

```
In [0]: # create inverse plots for the above models (predicting density with gain)

# transform data (take the sqrt(ln(#)) of y axis) and fit to linear regression
gauche$sqrt_ln_gain = sqrt(log(gauche$gain))
inv_reg_plot = ggplot(gauche, aes(x=sqrt_ln_gain, y=density)) +
  geom_point() +
  stat_smooth(method='lm', formula = y~x) +
  labs(title="ln Gain vs Density", y="Density", x="Natural Log of Gain")

# compute the residuals for the linear regression and plot
inv_lin_reg = lm(density ~ sqrt_ln_gain, data = gauche)
inv_residuals = data.frame(x=gauche$sqrt_ln_gain, resid=resid(inv_lin_reg))
inv_resid_plot = ggplot(inv_residuals, aes(x=x, y=resid)) +
  geom_point() +
  labs(title="ln Gain Residuals vs Density", x="ln(Gain)", y="Residuals")

# put both plots side by side
grid.arrange(inv_reg_plot, inv_resid_plot, ncol=2)

# print out the equation of the line and the sum of squared residuals
summarize_fit(inv_lin_reg)
```



```

In [0]: # create a confidence interval for the regression line

# find slope of regression line
slope = inv_lin_reg$coefficients[[2]]

# find standard deviation of slope
mse = sum(inv_lin_reg$residuals^2)/(nrow(gauge)-2)
deviations = var(gauge$sqrt_ln_gain)*(nrow(gauge)-1)
slope_sd = sqrt(mse/deviations)

# create confidence interval for slope
slope_me = qt(.975, nrow(gauge)-2, lower.tail = TRUE) * slope_sd
cat("95% Confidence Interval for slope: (", slope - slope_me, ",", slope
+ slope_me, ")\n")

# find y-int of regression line
yint = inv_lin_reg$coefficients[[1]]

# find standard deviation of y-int
meanx = mean(gauge$sqrt_ln_gain)
yint_sd = sqrt(mse*(1/nrow(gauge) + meanx))

# create confidence interval for y-intercept
yint_me = qt(.975, nrow(gauge)-2, lower.tail = TRUE) * yint_sd
cat("95% Confidence Interval for y-int: (", yint - yint_me, ",", yint + y
int_me, ")\n")

```

```

In [0]: # plot prediction interval with confidence bands

# plot data and regression line
newx <- seq(1.6, 2.5, .9/89)
plot(gauge$sqrt_ln_gain, gauge$density, ylim=c(0, .8), xlab="sqrt(log(Gai
n))", ylab="Density", main="Prediction Interval for Density", pch = 20)
abline(inv_lin_reg, col="lightblue")

# get upper and lower bounds for prediction interval of 90 points
conf_interval <- predict(inv_lin_reg, newdata=data.frame(sqrt_ln_gain = n
ewx), interval="prediction",
                        level = 0.95)
# plot points around regression line
lines(newx, conf_interval[,2], col="blue", lty=2)
lines(newx, conf_interval[,3], col="blue", lty=2)

```

```

In [0]: # create a function that produces a prediction interval for density, give
n a measured gain
pred_interval = function(input){
  transformed = sqrt(log(input))
  output = predict(lm(density ~ sqrt_ln_gain, data = gauge), newdata=da
ta.frame(sqrt_ln_gain = transformed), interval="prediction")
  return (c(output[2], output[3]))
}

```

```
In [0]: pred_interval(426.7)
```

```
In [0]: pred_interval(38.6)
```